

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import SGDClassifier
from sklearn.svm import SVC
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, ConfusionMatrixDisplay
import joblib
```

```
C:\Users\patil\anaconda3\lib\site-packages\pandas\core\computation\expressions.py:21: UserWarning: Pandas requires version '2.8.4' or newer of 'numexpr' (version '2.8.3' currently installed).
  from pandas.core.computation.check import NUMEXPR_INSTALLED
C:\Users\patil\anaconda3\lib\site-packages\pandas\core\arrays\masked.py:61: UserWarning: Pandas requires version '1.3.6' or newer of 'bottleneck' (version '1.3.5' currently installed).
  from pandas.core import (
```

```
In [2]: df = pd.read_csv(r'D:\Oasis\TASK 6\WineQT.csv')
df.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality	Id
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5	0
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5	1
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5	2
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6	3
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5	4

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1143 entries, 0 to 1142
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   fixed acidity    1143 non-null   float64
 1   volatile acidity 1143 non-null   float64
 2   citric acid      1143 non-null   float64
 3   residual sugar   1143 non-null   float64
 4   chlorides        1143 non-null   float64
 5   free sulfur dioxide 1143 non-null   float64
 6   total sulfur dioxide 1143 non-null   float64
 7   density          1143 non-null   float64
 8   pH               1143 non-null   float64
 9   sulphates        1143 non-null   float64
 10  alcohol          1143 non-null   float64
 11  quality          1143 non-null   int64  
 12  Id               1143 non-null   int64  
dtypes: float64(11), int64(2)
memory usage: 116.2 KB
```

In [4]: `df.describe()`

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	
count	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000
mean	8.311111	0.531339	0.268364	2.532152	0.086933	15.615486	45.914698	0.996730	3.311015	0.657708	1143.000000
std	1.747595	0.179633	0.196686	1.355917	0.047267	10.250486	32.782130	0.001925	0.156664	0.170399	1143.000000
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000	0.990070	2.740000	0.330000	1143.000000
25%	7.100000	0.392500	0.090000	1.900000	0.070000	7.000000	21.000000	0.995570	3.205000	0.550000	1143.000000
50%	7.900000	0.520000	0.250000	2.200000	0.079000	13.000000	37.000000	0.996680	3.310000	0.620000	1143.000000
75%	9.100000	0.640000	0.420000	2.600000	0.090000	21.000000	61.000000	0.997845	3.400000	0.730000	1143.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	68.000000	289.000000	1.003690	4.010000	2.000000	1143.000000

In [5]: `print("Missing value per column")`
`df.isna().sum()`

```
Missing value per column
Out[5]: fixed acidity      0
         volatile acidity    0
         citric acid        0
         residual sugar     0
         chlorides          0
         free sulfur dioxide 0
         total sulfur dioxide 0
         density            0
         pH                 0
         sulphates          0
         alcohol            0
         quality             0
         Id                 0
         dtype: int64
```

```
In [6]: if 'quality' not in df.columns:
    raise ValueError("Expected a column named 'quality'")
print("Quality distribution:")
print(df['quality'].value_counts().sort_index())

df['quality_binary'] = (df['quality'] >= 7).astype(int)

print("\nBinary label distribution:")
print(df['quality_binary'].value_counts())
```

```
Quality distribution:
quality
3      6
4     33
5    483
6    462
7   143
8    16
Name: count, dtype: int64
```


```
Binary label distribution:
quality_binary
0    984
1    159
Name: count, dtype: int64
```

```
In [7]: # Visualizations
sns.set_style("whitegrid")
```


```
plt.figure(figsize=(8,5))
sns.countplot(x='quality', data=df)
plt.title("Quality Distribution")
plt.show()


plt.figure(figsize=(8,5))
sns.countplot(x='quality_binary', data=df)
plt.title("Binary Quality Distribution")
plt.show()

# Correlation Heatmap
plt.figure(figsize=(12,10))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```



Binary Quality Distribution





fix
volat
resid
free sulfi
total sulfi
quali

```
In [8]: # Prepare data
FEATURE_COLS = [col for col in df.columns if col not in ['quality', 'quality_binary']]

X = df[FEATURE_COLS].values
y = df['quality'].values # multiclass target

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

print("Training set:", X_train.shape)
print("Testing set:", X_test.shape)
print("\nFeatures used:", FEATURE_COLS)
```

Training set: (914, 12)
Testing set: (229, 12)

Features used: ['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar', 'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density', 'pH', 'sulphates', 'alcohol', 'Id']

```
In [9]: # Classifier pipelines

# Random Forest
pipe_rf = Pipeline([
    ('scaler', StandardScaler()),
    ('rf', RandomForestClassifier(random_state=42))])
# SGDClassifier
pipe_sgd = Pipeline([
    ('scaler', StandardScaler()),
    ('sgd', SGDClassifier(max_iter=5000, tol=1e-3, random_state=42))])

# SVC
pipe_svc = Pipeline([
    ('scaler', StandardScaler()),
    ('svc', SVC(probability=True, random_state=42))])

models = {
    "RandomForest": pipe_rf,
    "SGDClassifier": pipe_sgd,
    "SVC": pipe_svc
}
```

```
In [10]: # 8. Train & evaluate
```

```
results = {}

for name, model in models.items():
    print(f"\n--- Training {name} ---")
    model.fit(X_train, y_train)

    preds = model.predict(X_test)
    acc = accuracy_score(y_test, preds)
    results[name] = acc

    print(f"{name} Accuracy: {acc:.4f}")
    print(classification_report(y_test, preds))


# Confusion Matrix
plt.figure(figsize=(6,4))
ConfusionMatrixDisplay(confusion_matrix(y_test, preds)).plot(cmap="Blues")
plt.title(f"{name} - Confusion Matrix")
plt.show()
```

```
--- Training RandomForest ---
```

```
RandomForest Accuracy: 0.7162
```

	precision	recall	f1-score	support
3	0.00	0.00	0.00	1
4	0.00	0.00	0.00	7
5	0.73	0.85	0.78	97
6	0.68	0.73	0.71	92
7	0.79	0.52	0.62	29
8	0.00	0.00	0.00	3
accuracy			0.72	229
macro avg	0.37	0.35	0.35	229
weighted avg	0.68	0.72	0.69	229

```
C:\Users\patil\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, msg_start, len(result))  
C:\Users\patil\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, msg_start, len(result))  
C:\Users\patil\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, msg_start, len(result))  
<Figure size 600x400 with 0 Axes>
```



```
--- Training SGDClassifier ---
SGDClassifier Accuracy: 0.5677
      precision    recall  f1-score   support

       3         0.00     0.00     0.00        1
       4         0.00     0.00     0.00        7
       5         0.67     0.71     0.69       97
       6         0.58     0.45     0.50       92
       7         0.37     0.69     0.48       29
       8         0.00     0.00     0.00        3

   accuracy                           0.57      229
macro avg       0.27     0.31     0.28      229
weighted avg    0.56     0.57     0.56      229
```

```
C:\Users\patil\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
C:\Users\patil\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```


```
_warn_prf(average, modifier, msg_start, len(result))
```

```
C:\Users\patil\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
<Figure size 600x400 with 0 Axes>
```

SGDClassifier - Confusion Matrix




--- Training SVC ---

SVC Accuracy: 0.6725

	precision	recall	f1-score	support
3	0.00	0.00	0.00	1
4	0.00	0.00	0.00	7
5	0.71	0.82	0.76	97
6	0.63	0.70	0.66	92
7	0.71	0.34	0.47	29
8	0.00	0.00	0.00	3
accuracy			0.67	229
macro avg	0.34	0.31	0.31	229
weighted avg	0.64	0.67	0.65	229

```
C:\Users\patil\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, msg_start, len(result))  
C:\Users\patil\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, msg_start, len(result))  
C:\Users\patil\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, msg_start, len(result))  
<Figure size 600x400 with 0 Axes>
```



```
In [11]: # 9. Cross-validation  
for name, model in models.items():  
    print(f"\nCross-validating {name}...")
```

```
scores = cross_val_score(model, X, y, cv=5, scoring='accuracy')
print(f"{name} CV Mean Accuracy: {scores.mean():.4f}")
```

```
Cross-validating RandomForest...
RandomForest CV Mean Accuracy: 0.5425

Cross-validating SGDClassifier...
SGDClassifier CV Mean Accuracy: 0.5372

Cross-validating SVC...
SVC CV Mean Accuracy: 0.5835
```

```
In [12]: # 10. Hyperparameter tuning
```

```
# Random Forest Grid
rf_grid = {
    'rf__n_estimators': [100, 200],
    'rf__max_depth': [None, 10, 20]
}

rf_search = GridSearchCV(pipe_rf, rf_grid, cv=4, scoring='accuracy', n_jobs=-1)
rf_search.fit(X_train, y_train)

print("Best RF Params:", rf_search.best_params_)
print("Best RF Accuracy:", rf_search.best_score_)

# SVC Grid
svc_grid = {
    'svc__C': [0.1, 1, 10],
    'svc__kernel': ['rbf', 'linear']
}

svc_search = GridSearchCV(pipe_svc, svc_grid, cv=4, scoring='accuracy', n_jobs=-1)
svc_search.fit(X_train, y_train)

print("Best SVC Params:", svc_search.best_params_)
print("Best SVC Accuracy:", svc_search.best_score_)
```

```
Best RF Params: {'rf__max_depth': None, 'rf__n_estimators': 200}
Best RF Accuracy: 0.6444543399984678
Best SVC Params: {'svc__C': 1, 'svc__kernel': 'rbf'}
Best SVC Accuracy: 0.6094767486401593
```

```
In [13]: # 11. Save models
```

```
joblib.dump(rf_search.best_estimator_, "best_random_forest.joblib")
joblib.dump(svc_search.best_estimator_, "best_svc.joblib")
```

```
joblib.dump(pipe_sgd, "baseline_sgd.joblib")
print("Models saved successfully.")
```

Models saved successfully.

```
In [14]: # 12. Example prediction
sample = X_test[0].reshape(1, -1)

print("Sample Prediction:")
print("RF:", rf_search.best_estimator_.predict(sample)[0])
print("SVC:", svc_search.best_estimator_.predict(sample)[0])
print("SGD:", pipe_sgd.predict(sample)[0])
```

Sample Prediction:
RF: 6
SVC: 6
SGD: 6

```
In [ ]:
```