

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, RobustScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, roc_auc_score, RocCurveDisplay

from sklearn.ensemble import IsolationForest
import joblib
```

```
C:\Users\patil\anaconda3\lib\site-packages\pandas\core\computation\expressions.py:21: UserWarning: Pandas requires version '2.8.4' or newer of 'numexpr' (version '2.8.3' currently installed).
  from pandas.core.computation.check import NUMEXPR_INSTALLED
C:\Users\patil\anaconda3\lib\site-packages\pandas\core\arrays\masked.py:61: UserWarning: Pandas requires version '1.3.6' or newer of 'bottleneck' (version '1.3.5' currently installed).
  from pandas.core import (
```

```
In [4]: df = pd.read_csv(r'D:\Oasis\TASK 7\archive (11)\creditcard.csv')
df.head()
```

```
Out[4]:   Time      V1      V2      V3      V4      V5      V6      V7      V8      V9 ...      V21      V22      V23
  0    0.0 -1.359807 -0.072781  2.536347  1.378155 -0.338321  0.462388  0.239599  0.098698  0.363787 ... -0.018307  0.277838 -0.110474
  1    0.0  1.191857  0.266151  0.166480  0.448154  0.060018 -0.082361 -0.078803  0.085102 -0.255425 ... -0.225775 -0.638672  0.101288
  2    1.0 -1.358354 -1.340163  1.773209  0.379780 -0.503198  1.800499  0.791461  0.247676 -1.514654 ...  0.247998  0.771679  0.909412
  3    1.0 -0.966272 -0.185226  1.792993 -0.863291 -0.010309  1.247203  0.237609  0.377436 -1.387024 ... -0.108300  0.005274 -0.190321
  4    2.0 -1.158233  0.877737  1.548718  0.403034 -0.407193  0.095921  0.592941 -0.270533  0.817739 ... -0.009431  0.798278 -0.137458
```

5 rows × 31 columns

```
In [5]: df.info()
print("\nStatistical Summary:")
```

```
df.describe()  
print("\nMissing Values:")  
df.isna().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column   Non-Null Count   Dtype  
 --- 
 0   Time     284807 non-null   float64
 1   V1       284807 non-null   float64
 2   V2       284807 non-null   float64
 3   V3       284807 non-null   float64
 4   V4       284807 non-null   float64
 5   V5       284807 non-null   float64
 6   V6       284807 non-null   float64
 7   V7       284807 non-null   float64
 8   V8       284807 non-null   float64
 9   V9       284807 non-null   float64
 10  V10      284807 non-null   float64
 11  V11      284807 non-null   float64
 12  V12      284807 non-null   float64
 13  V13      284807 non-null   float64
 14  V14      284807 non-null   float64
 15  V15      284807 non-null   float64
 16  V16      284807 non-null   float64
 17  V17      284807 non-null   float64
 18  V18      284807 non-null   float64
 19  V19      284807 non-null   float64
 20  V20      284807 non-null   float64
 21  V21      284807 non-null   float64
 22  V22      284807 non-null   float64
 23  V23      284807 non-null   float64
 24  V24      284807 non-null   float64
 25  V25      284807 non-null   float64
 26  V26      284807 non-null   float64
 27  V27      284807 non-null   float64
 28  V28      284807 non-null   float64
 29  Amount    284807 non-null   float64
 30  Class     284807 non-null   int64  
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

Statistical Summary:

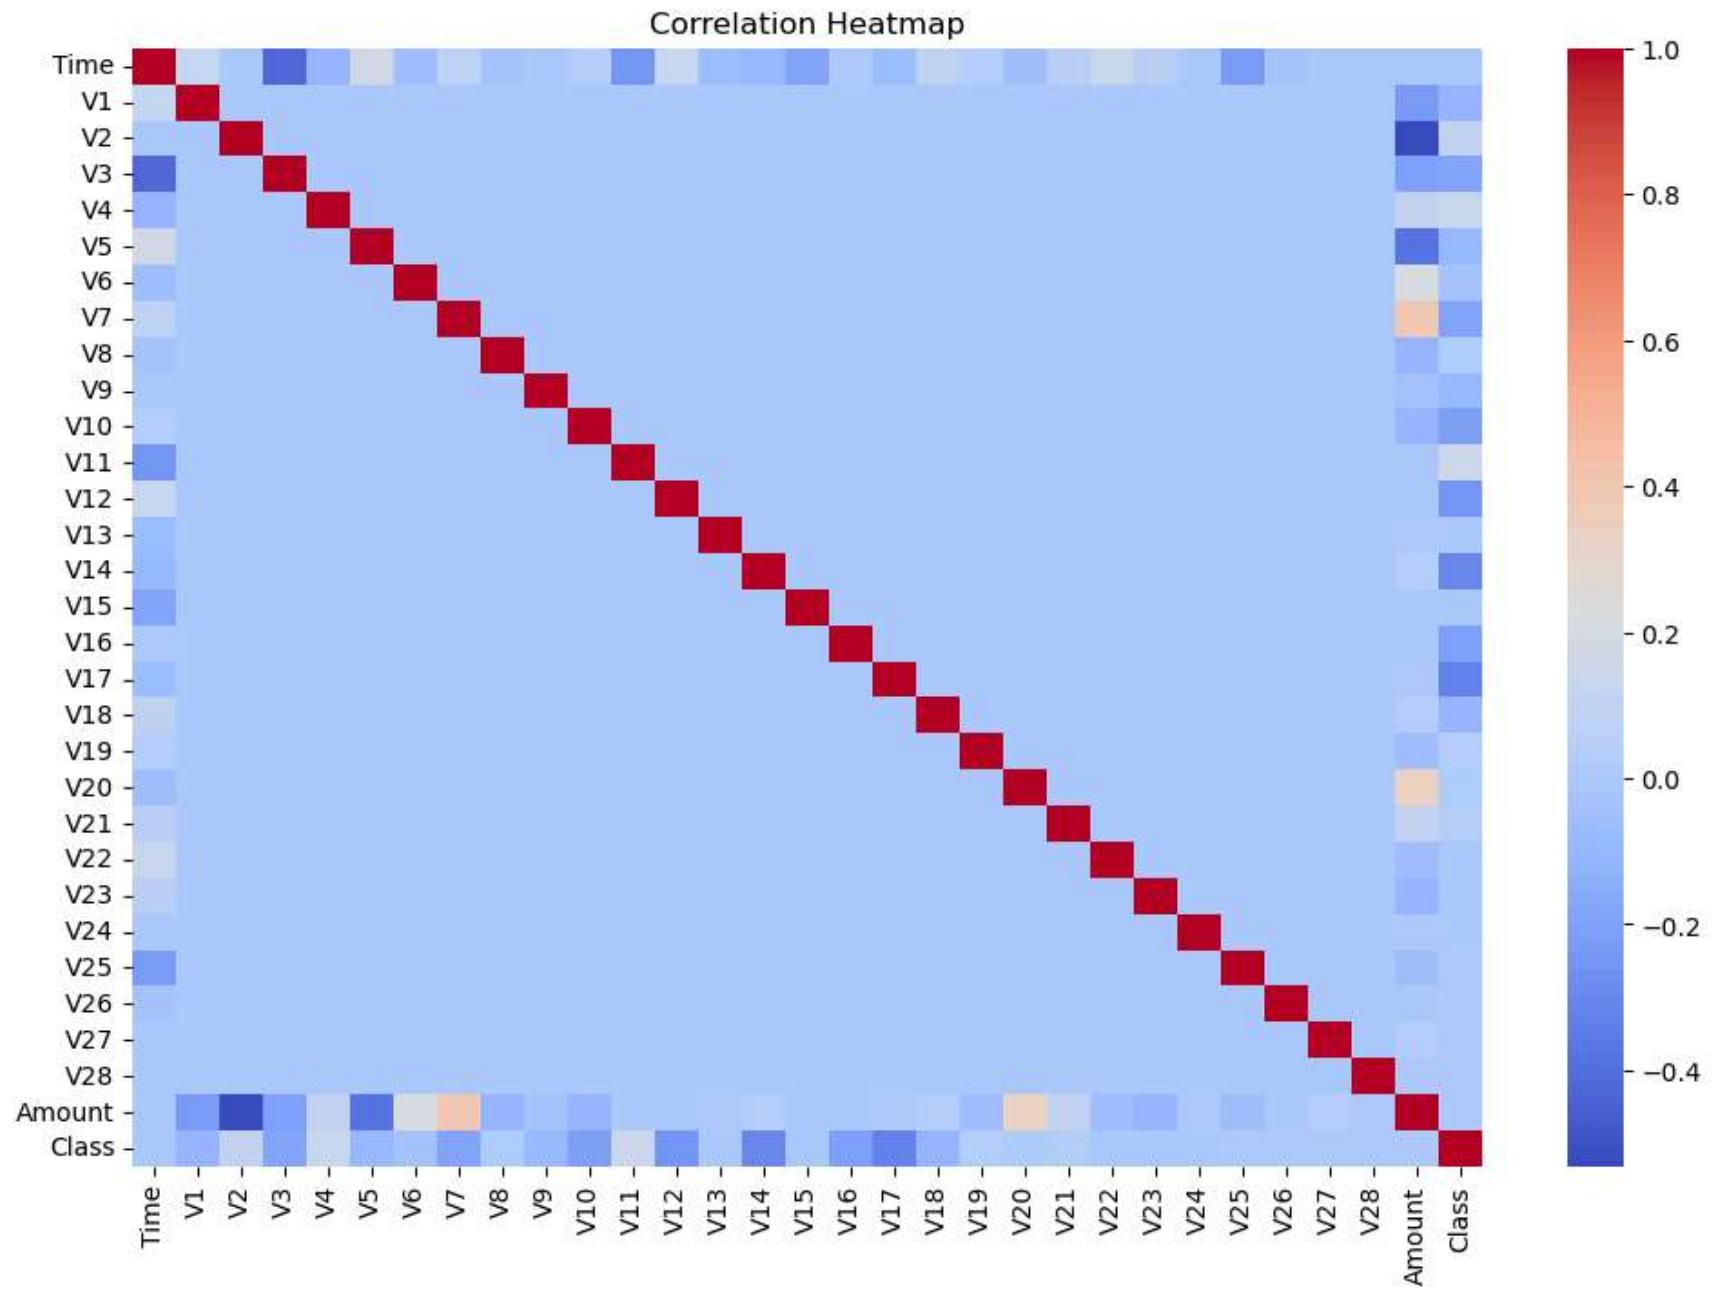
Missing Values:

```
Out[5]: Time      0  
          V1       0  
          V2       0  
          V3       0  
          V4       0  
          V5       0  
          V6       0  
          V7       0  
          V8       0  
          V9       0  
          V10      0  
          V11      0  
          V12      0  
          V13      0  
          V14      0  
          V15      0  
          V16      0  
          V17      0  
          V18      0  
          V19      0  
          V20      0  
          V21      0  
          V22      0  
          V23      0  
          V24      0  
          V25      0  
          V26      0  
          V27      0  
          V28      0  
          Amount    0  
          Class    0  
          dtype: int64
```

```
In [8]: #Fraud Distribution  
print(df['Class'].value_counts())  
sns.countplot(x='Class', data=df)  
plt.title("Fraud vs Non-Fraud Distribution")  
plt.show()  
  
Class  
0    284315  
1     492  
Name: count, dtype: int64
```



```
In [9]: plt.figure(figsize=(12,8))
sns.heatmap(df.corr(), cmap='coolwarm', annot=False)
plt.title("Correlation Heatmap")
plt.show()
```



```
In [10]: # Split features and target
X = df.drop("Class", axis=1)
y = df["Class"]
```

```
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
scaler = RobustScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [11]: models = {
    "Logistic Regression": LogisticRegression(max_iter=5000, class_weight='balanced'),
    "Decision Tree": DecisionTreeClassifier(class_weight='balanced', random_state=42),
    "Neural Network": MLPClassifier(hidden_layer_sizes=(64,32), max_iter=300)
}
```

```
In [12]: for name, model in models.items():
    print(f"\nTraining {name}...")
    model.fit(X_train_scaled, y_train)
    preds = model.predict(X_test_scaled)

    print(f"{name} Classification Report:")
    print(classification_report(y_test, preds))

    # AUC Score
    auc = roc_auc_score(y_test, preds)
    print("AUC Score:", auc)

    # Confusion Matrix
    cm = confusion_matrix(y_test, preds)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title(f"{name} Confusion Matrix")
    plt.show()
```

Training Logistic Regression...

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	1.00	0.98	0.99	56864
1	0.06	0.92	0.11	98
accuracy			0.98	56962
macro avg	0.53	0.95	0.55	56962
weighted avg	1.00	0.98	0.99	56962

AUC Score: 0.9469791152251559

Logistic Regression Confusion Matrix



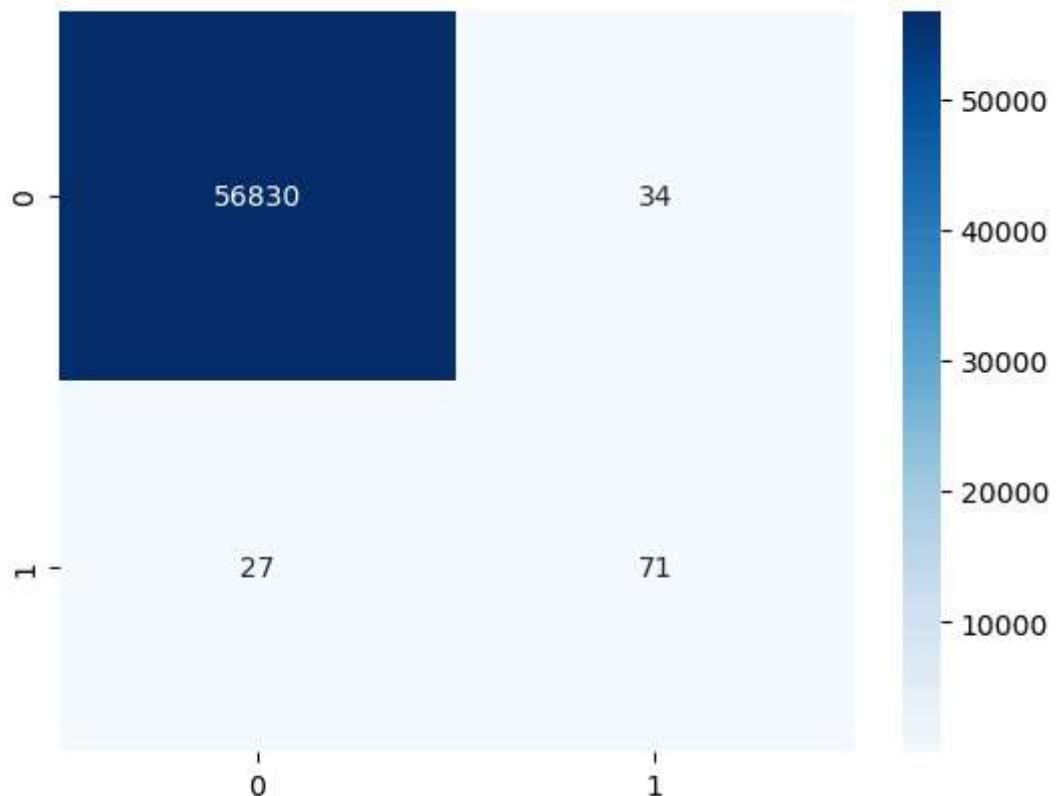
Training Decision Tree...

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.68	0.72	0.70	98
accuracy			1.00	56962
macro avg	0.84	0.86	0.85	56962
weighted avg	1.00	1.00	1.00	56962

AUC Score: 0.8619459390396564

Decision Tree Confusion Matrix

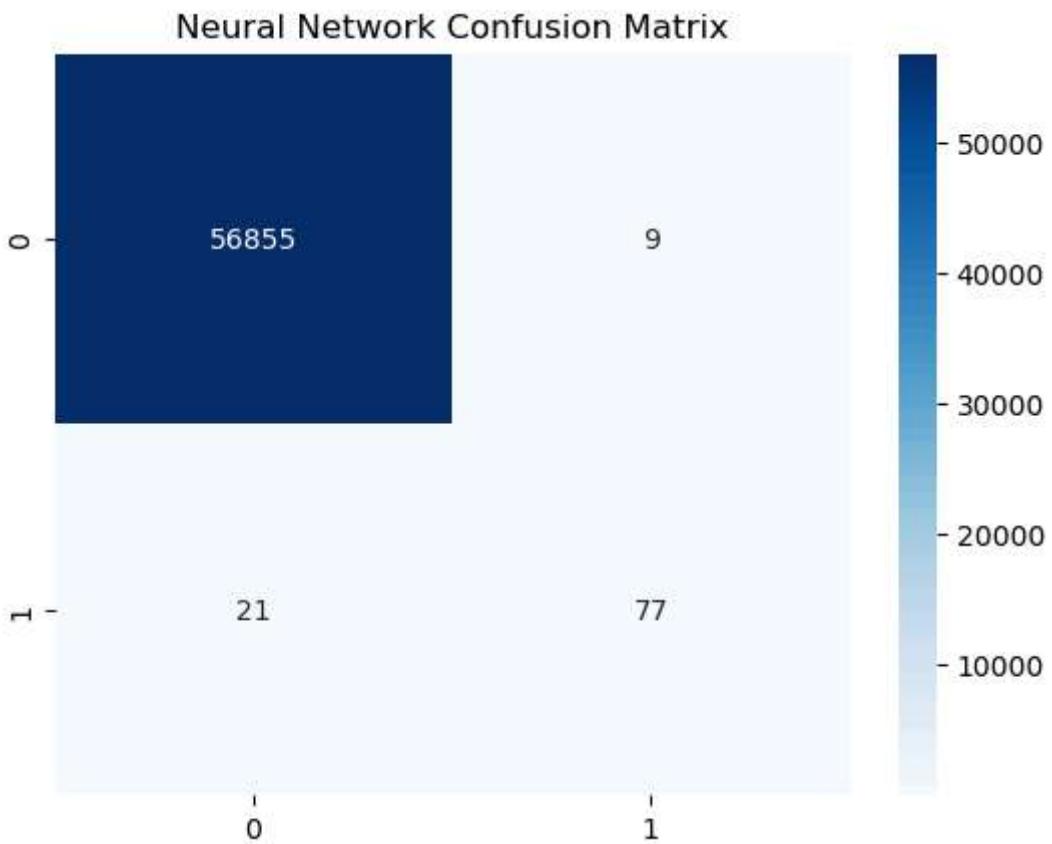


Training Neural Network...

Neural Network Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.90	0.79	0.84	98
accuracy			1.00	56962
macro avg	0.95	0.89	0.92	56962
weighted avg	1.00	1.00	1.00	56962

AUC Score: 0.8927780066725619



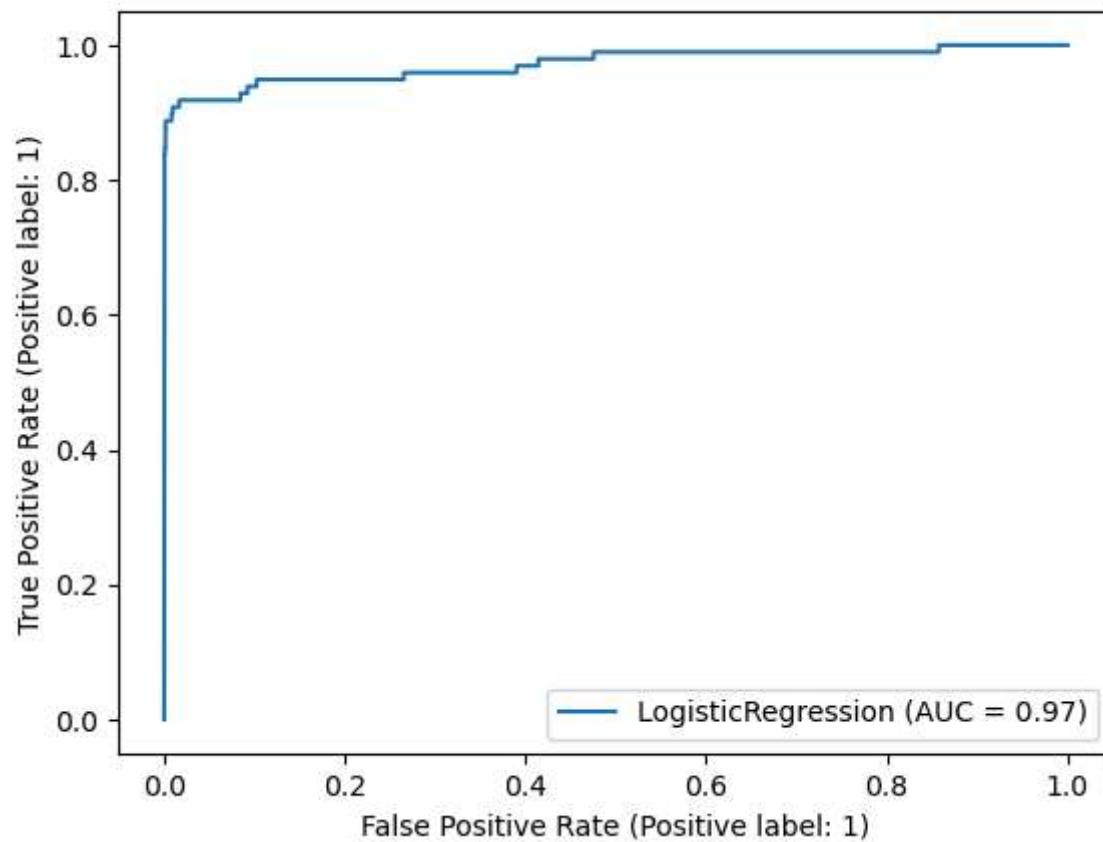
```
In [13]: plt.figure(figsize=(8,6))

for name, model in models.items():
    RocCurveDisplay.from_estimator(model, X_test_scaled, y_test)
    plt.title("ROC Curves for Models")

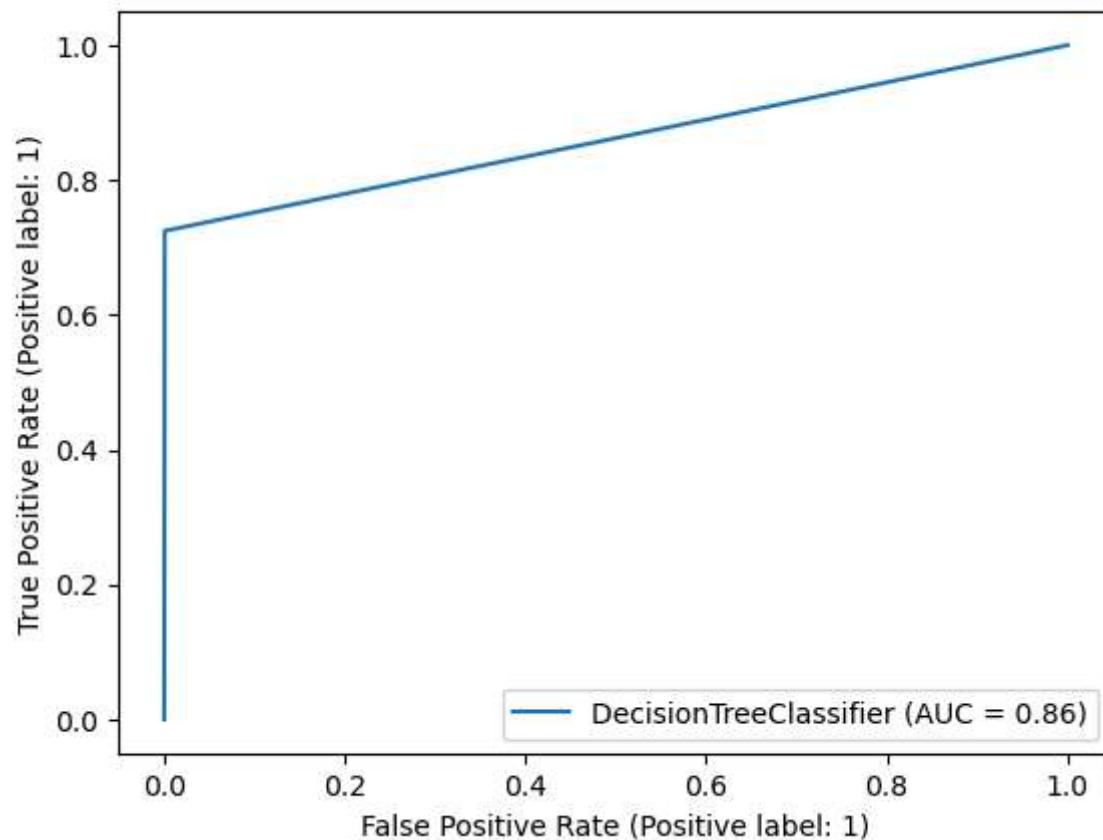
plt.show()
```

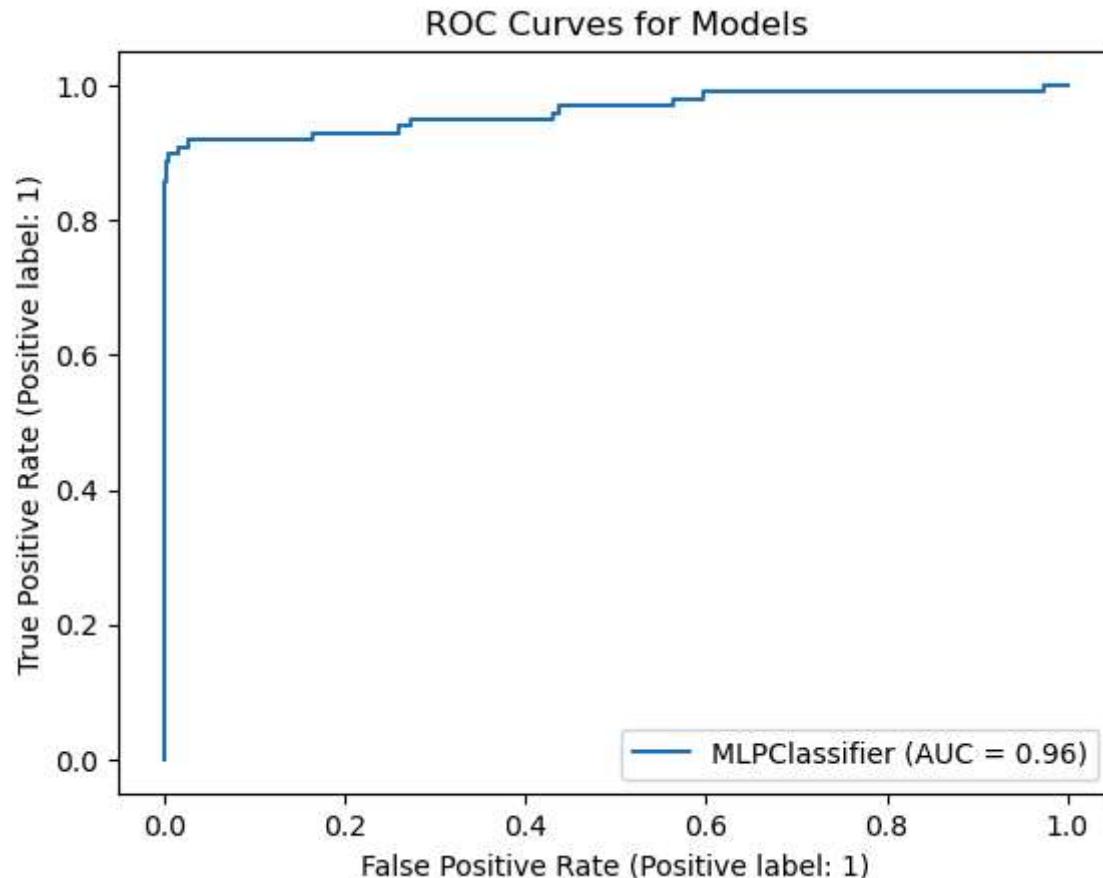
<Figure size 800x600 with 0 Axes>

ROC Curves for Models



ROC Curves for Models





```
In [14]: iso = IsolationForest(contamination=0.001, random_state=42)
iso.fit(X_train_scaled)
anomaly_scores = iso.predict(X_test_scaled)
anomaly_scores = [1 if x == -1 else 0 for x in anomaly_scores]
print("Isolation Forest Report:")
print(classification_report(y_test, anomaly_scores))
```

```
Isolation Forest Report:
      precision    recall   f1-score   support

          0       1.00     1.00     1.00    56864
          1       0.39     0.26     0.31      98

   accuracy                           1.00    56962
macro avg       0.69     0.63     0.65    56962
weighted avg    1.00     1.00     1.00    56962
```

```
In [17]: param_grid = {
    'max_depth': [4, 6, 8, None],
    'min_samples_split': [2, 5, 10]}

grid = GridSearchCV(
    DecisionTreeClassifier(class_weight='balanced'),
    param_grid,
    scoring='f1',
    cv=3)

grid.fit(X_train_scaled, y_train)
print("Best Params:", grid.best_params_)
print("Best Score:", grid.best_score_)
```

Best Params: {'max_depth': None, 'min_samples_split': 2}
 Best Score: 0.7295246840701387

```
In [18]: joblib.dump(models["Logistic Regression"], "logistic_fraud.joblib")
joblib.dump(models["Decision Tree"], "decisiontree_fraud.joblib")
joblib.dump(models["Neural Network"], "nn_fraud.joblib")

print("Models saved successfully.")
```

Models saved successfully.

```
In [19]: sample = X_test_scaled[0].reshape(1, -1)

print("Logistic Prediction:", models["Logistic Regression"].predict(sample))
print("Decision Tree Prediction:", models["Decision Tree"].predict(sample))
print("Neural Net Prediction:", models["Neural Network"].predict(sample))
```

Logistic Prediction: [0]
 Decision Tree Prediction: [0]
 Neural Net Prediction: [0]

